

All-agent analysis of recent games against strong opponents

Reconciled result from independent replication and two-way cross-review

Current bot agent 6589709, submission 41079653, source SHA-256 6f992a5a...
Evidence 153 new open games; one direct rank-13 game; 2,684 top-20-source benchmark games
Reviewers local_codex_1, claude_1, chatgpt_1
Frozen package commit 73718b3fdf9f2dc13359e17cb0ce002f95ea559e

Final decision. The review leaves one research direction: H3a, which increases priority for eligible opponent-created trees only after the public state shows at least three opponent workers. It is **not ready to run**: the conditioned source, exact value runner, byte-equivalence bridge, and required multi-game trajectory package do not exist. No second or third improvement survived cross-review.

1. Evidence boundary

The package contains 153 current-bot games: **95 wins, 2 ties and 56 losses**, with 68 games in seat 0 and 85 in seat 1. Seven sealed-tagged games were excluded. The direct strong-opponent evidence is only game 897780884 against rank-13 Astrobytes; the larger top-20 corpus is observational context, not a direct treatment comparison.

Evidence slice	Size	Permitted interpretation
Current-new open games	153	Current behavior, associations and recurrent failure modes
Full 300-turn current games	96	Phase-by-phase score decomposition
Direct current vs top-20	1	Exact postmortem only; not broad evidence
Top-20-source open games	2,684	Benchmark context; not causal value
Sealed-tagged games	7	Excluded and unopened

No report used sealed data, performed a build or simulation, created a candidate, or mutated the Arena.

2. The recurrent loss pattern

The blind replication isolated ten catastrophic games, defined before mechanism analysis as final margin at most -100 :

- **10/153 games** (6.54%; Wilson 95% interval 3.59–11.61%);
- total margin $-1,674$, or -167.4 per catastrophe;
- every catastrophic opponent finished with 3–4 workers while the resident stayed at exactly two;
- the resident was still ahead at turn 150 in 9/10, but ahead at turn 250 in only 1/10 and at the end in 0/10.

Checkpoint	t50	t100	t150	t200	t250	Final
Resident ahead, catastrophes	9/10	10/10	9/10	6/10	1/10	0/10

This is a late crossover under opponent scaling, not a recurrent opening collapse. It is still observational: final roster is endogenous, the resident has no workforce variation, and opponent scaling bundles many policy differences.

3. What changes across the game

Claude's non-blind replication independently decomposed points per 50-turn window over all 96 full games. Against opponents that finish with at least three workers, resident scoring does not decline; opponent scoring accelerates sharply.

Window	Our points	Their points	Net	Opponent roster group
t50-100	36.92	8.64	+28.28	≥ 3 workers
t100-150	37.64	16.47	+21.17	≥ 3 workers
t150-200	39.75	46.06	-6.31	≥ 3 workers
t200-250	40.83	79.81	-38.97	≥ 3 workers
t250-300	42.56	102.19	-59.64	≥ 3 workers

For the 60 full games against opponents finishing with at most two workers, the resident's net remains positive in every window. For the 36 games against 3-4-worker opponents, the net turns negative in t150-200 and deteriorates afterward. The reconciled interpretation is deliberately narrow:

Opponent workforce divergence can be established by turn 150; the score crossover and largest terminal damage occur later.

The decomposition rejects the simple story that our own output collapses. It does not, by itself, prove every marginal own-economy improvement worthless; those families remain closed by their separate controlled experiments and displacement accounting.

4. Direct game: Astrobytes, 897780884

The resident lost 333-403. It trained once at turn 11; Astrobytes trained at turns 1, 56 and 105 and finished with four workers. The resident led through turn 200, then crossed behind during the late acceleration. The direct trajectory supports the chronology but cannot identify a treatment effect.

The early WAIT sequence at turns 4-8 is followed by harvesting the same tree and matches the documented ripening-wait behavior. Across the 36 scaled-opponent full games, resident WAIT count correlates only -0.046 with margin. Cross-review therefore rejected a generic WAIT legality/precedence audit as a ranked improvement.

Tree lineage, first feller, wood recipient, and counterfactual legality are unavailable for the wider cohort in this package. Claims depending on those fields remain unavailable.

5. What each agent contributed

Agent	Principal contribution	Correction after review
local_codex_1	Initial integrated ranking from quantitative, action-flow and direct-postmortem tracks; H3a first.	Original ranks 2–3 do not survive the full replication ring.
claude_1	96-game score-window decomposition: opponent acceleration, not resident output decline. Prior exposure disclosed.	Withdrew the endgame removal-race idea and kept rank 3 empty.
chatgpt_1	Blind ten-catastrophe study, exact IDs/arithmetic, late-crossover chronology, and H3a trigger preflight.	Folded its discriminator into H3a and withdrew the WAIT audit.

Both cross-reviews returned **ACCEPT WITH CORRECTIONS**. Claude reproduced every checkable number in the blind report. ChatGPT accepted the phase decomposition but narrowed its causal wording. This agreement is stronger than the initial ranking because it survived independent recomputation and adversarial review.

6. Final ranked disposition

Rank	Idea	Disposition	Reason
1	H3a: pressure-conditioned opponent-crop priority	BLOCKED, TOP FUTURE ROUTE	Unanimous direction and exact always-on reconstruction, but no conditioned source, value runner, equality bridge, or runnable multi-game preflight package.
—	Endgame conversion removal race	REJECTED	One direct game; required tree/feller/wood attribution absent; correlation is collinear with opponent roster; no defensible cohort value.
—	B3.14 sticky banking	MEASUREMENT ONLY	17 unbanked wood imply at most +0.444 own points per game; too small for an improvement rank.
—	WAIT legal-ity/precedence	REJECTED	Same-tree ripening behavior, no cohort signal, legality counterfactual unavailable.

There is intentionally no rank 2 or rank 3. Padding the list with low-value audits would hide the central result: the analysis found one scientifically live direction, and it is not yet executable.

7. H3a: exact next decision

The intended three arms remain:

1. unchanged resident control;
2. exact eligible opponent-crop priority always on;
3. the identical priority activated, and then kept active, after public state first shows at least three opponent workers.

The existing reconstruction self-test passes, but it proves source reconstruction only. Before any 6,144-task value panel, a new host-authorized read-only task must materialize the required trajectories and test the trigger gate: sufficiently early activation in the ten catastrophes, low false activation on matched wins, and eligible $\text{ETA} \leq 6$ decisions after activation. If that preflight fails, stop H3a without tuning it. If it passes, build the conditioned source and exact value runner under a new frozen protocol; the frozen protocol's value gates govern.

A purely illustrative 20% recovery of the ten-catastrophe negative mass would be $1,674 \times 0.20/153 = +2.19$ mean margin per current game. This is a scenario, not a prediction and not an Arena-rating conversion.

8. Data and provenance corrections

- Do not compute `planted_ok / plant_cmd`: over top-20 sides the aggregate numerator is 86,023 and the denominator 81,280. The counters are not documented as a subset relation, so a success rate can exceed 100% and is unsound.
- For retrospective “scaled by turn 150” counts, require a successfully landed extra worker:

```
second_train_turn <= 151
AND roster_final >= 3
```

Game 897782434 contains a failed TRAIN and must be excluded.

- Do not use the earlier unexplained “1,268” count; no coherent package predicate reproduced it.
- The rubric's SHA-256 and Git blob hash are different hash types over the same content, not a provenance failure.

9. Source artifacts

- Local synthesis: `local_codex_1/top-player-new-games-final-ranked-ideas-2026-08-02.md`
- Claude replication: `claude_1/top-player-full-review-replication-2026-08-02.md`
- Claude cross-review: `claude_1/top-player-full-review-cross-review-of-chatgpt_1-2026-08-02.md`
- ChatGPT blind replication: `chatgpt_1/top-player-full-review-replication-2026-08-02.md`
- ChatGPT cross-review: `chatgpt_1/top-player-full-review-cross-review-of-claude_1-2026-08-02.md`